



Contract #: DTFH6116D00035

TOPR #: HOIT212116217

JPO

Operational Data Environment

Interim Design Document

Submitted to:

U.S. Department of Transportation (USDOT)

Federal Highway Administration ITS JPO

January 16, 2017

Prepared by:

Booz Allen Hamilton
8283 Greensboro Drive
McLean, VA 22102

This page intentionally blank

CONTENTS

Version History.....	4
1 Introduction	5
2 Project Overview.....	5
3 System Overview	6
3.1 Producer Mechanisms	6
3.2 Consumer Mechanisms	7
3.3 ODE Management Console	7
4 Architecture Pattern	8
4.1 Pattern Description	8
4.2 Pattern Topology	9
5 JPO ODE Micro-services Topology.....	12
5.1 Deployments.....	13
6 Appendix.....	14
6.1 Glossary	14

VERSION HISTORY

Version #	Implemented By	Revision Date	What Changed?
1.0	Hamid Musavi Tony Chen	1/11/2017	Initial release of the preliminary/interim design

1 INTRODUCTION

The JPO Operational Data Environment (ODE) product is being developed under Agile Development Methodologies, using an open architecture approach, in an open source environment. This document describes the preliminary architectural design of the JPO ODE and its interfaces with external systems including the TMC applications, field devices and center services.

All stakeholders are invited to provide input to this document. Stakeholders should direct all input on this document to the JPO Product Owner at DOT, FHWA, JPO.

This document is a living document and will be updated throughout the life of the JPO ODE project to reflect the most recent changes in the ODE design and stakeholder feedback.

2 PROJECT OVERVIEW

An Operational Data Environment is a real-time data acquisition and distribution software system that processes and routes data from Connected-X devices – including connected vehicles (CV), personal mobile devices, infrastructure components, and sensors – to subscribing applications to support the operation, maintenance, and use of the transportation system, as well as related research and development efforts.

The ODE is intended to complement a connected vehicle infrastructure by brokering, processing and routing data from various data sources, including connected vehicles, field devices, Transportation Management Center (TMC) applications and a variety of other data users. Data users include but not limited to transportation software applications, Research Data Exchange (RDE), US DOT Situation Data Warehouse.

As a data provisioning service, the ODE can provision data from disparate data sources to software applications that have placed data subscription requests to the ODE. On the other direction, the ODE can accept data from CV applications and broadcast them to field devices through Road Side Units (RSU) and US DOT Situation Data Warehouse which in turn will transmit the data to Sirius XM satellites for delivery to the connected vehicles in the field.

While provisioning data from data sources to data users, the ODE also will perform necessary security / credential checks and, as needed, data validation and sanitization.

- Data validation is the process of making a judgment about the quality of the data and handling invalid data as prescribed by the system owners.
- Data sanitization is the modification of data as originally received to reduce or eliminate the possibility that the data can be used to compromise the privacy of the individual(s) that might be linked to the data.

3 SYSTEM OVERVIEW

JPO ODE is an open-sourced software application that will enable the transfer of data between field devices and backend TMC systems for operational, monitoring, and research purposes. The system will enable applications to submit data through a variety standard interfaces as illustrated in the figure below.

The mechanisms chosen for a specific deployment will depend on the infrastructure, technical resources, and applications available to an ODE environment.

The JPO-ODE will be designed to support the producers and consumers of CV data as illustrated in Figure 1 below. ***The implementation timeline for the identified interfaces will depend on the needs of the JPO ODE customers (Wyoming CV Pilot site, initially) and the priority of these capabilities to the JPO-ODE product owner.***

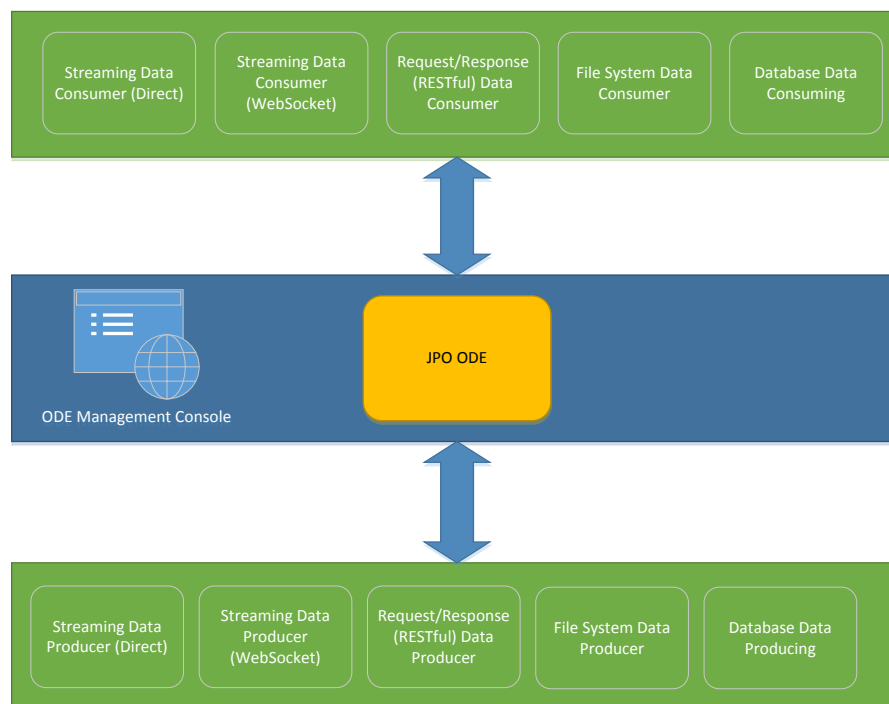


Figure 1 - ODE System Data Producers and Consumers

3.1 PRODUCER MECHANISMS

The JPO-ODE will be designed to support the following mechanisms for inputting ASN.1 encoded BSMs, TIM messages in a human readable encoded format (e.g. JSON), environmental and various other system logs.

- **Streaming Data Producers (Direct):** Applications can directly interact with the messaging service through the use of the service's native API and publish messages to

be processed by the ODE. *This interface will be available only to applications residing inside a private network domain.*

- **Streaming Data Producers (WebSocket):** Applications can interact with the messaging service and publish messages to be processed by the ODE. *This interface will be available to all applications whether residing in the private network domain or in the cloud. For cloud applications Secure WebSocket (wss) protocol will be required.*
- **RESTful API Data Producers:** Applications can connect with the ODE through a RESTful API and submit messages to the messaging service through HTTP POST commands. *This interface will be available to all applications whether residing in the private network domain or in the cloud. For cloud applications Secure HTTP (https) protocol will be required.*
- **File System Data Producers:** Encoded message files and log files messages can be dropped into a shared file system location and systematically pulled in to the data broker. *This interface will be available to applications residing in the private network domain or in the cloud. This interface will only available through Secure Copy (scp).*
- **Database Data Producer:** A shared database where encoded messages are stored can also be connected directly into the ODE to monitor and process new records. *This interface will be available only to applications residing in the private network domain.*

3.2 CONSUMER MECHANISMS

The JPO-ODE will be designed to support the following mechanisms for outputting decoded BSM, Map and Spat data. ¹

- **Streaming Data Consumers (Direct):** Applications can subscribe directly to the messaging service through the use of the messaging service's native API. *This interface will be available only to applications residing in the private network domain.*
- **Streaming Data Consumers (WebSocket):** Applications can subscribe to the messaging service through the use of a standard WebSocket API. *This interface will be available to all applications whether residing in the private network domain or in the cloud. For cloud applications Secure WebSocket (wss) protocol will be required.*
- **RESTful API Data Consumers:** Applications can connect directly with a RESTful API and submit messages to the messaging service through HTTP commands. *This interface will be available to all applications whether residing in the private network domain or in the cloud. For cloud applications Secure HTTP (https) protocol will be required.*
- **File System Data Consumers:** Through the use of a shared file repository, applications can monitor collection of data messages. *This interface will be available to applications residing in the private network domain or in the cloud. This interface will only available through Secure Copy (scp).*
- **Database Data Consumers:** Data messages can be directly inserted into a shared application database and made available for queries.

3.3 ODE MANAGEMENT CONSOLE

¹ Initial release of the ODE will only support BSM. Map and Spat will be supported in follow on releases as other CV pilots will adopt the JPO-ODE for their deployment.

Ode will provide a management console for a role-based Identity and Access Management and for SNMPP device management and provisioning.

4 ARCHITECTURE PATTERN

JPO ODE will be developed according to the micro-services architecture pattern. The micro-services architecture pattern is a highly scalable design pattern and a viable alternative to monolithic applications² and service-oriented architectures.

4.1 PATTERN DESCRIPTION

The micro-services pattern consists of three major concepts:

1. *Separately deployed units*: As illustrated in Figure 2, each component of the micro-services architecture is deployed as a separate unit, allowing for easy deployment, increased scalability, and a high degree of component decoupling.
2. *Service component*: In micro-services architecture, we deal with service components, which can vary in granularity from a single module to a large portion of the application. Service components contain one or more modules (Java classes) that represent either a single-purpose function (e.g., decode BSMs from ASN.1) or an independent portion of a large business application (e.g., sanitize BSM data according to the client request).
3. *Distributed architecture*: All the components within the architecture are fully decoupled from one other and accessed through a messaging service. This concept is what allows microservices architecture pattern achieve some of its superior scalability and deployment characteristics.

² In software engineering, a **monolithic application** describes a single-tiered software **application** in which the user interface and data access code are combined into a single program from a single platform. A **monolithic application** is self-contained, and independent from other computing **applications**.

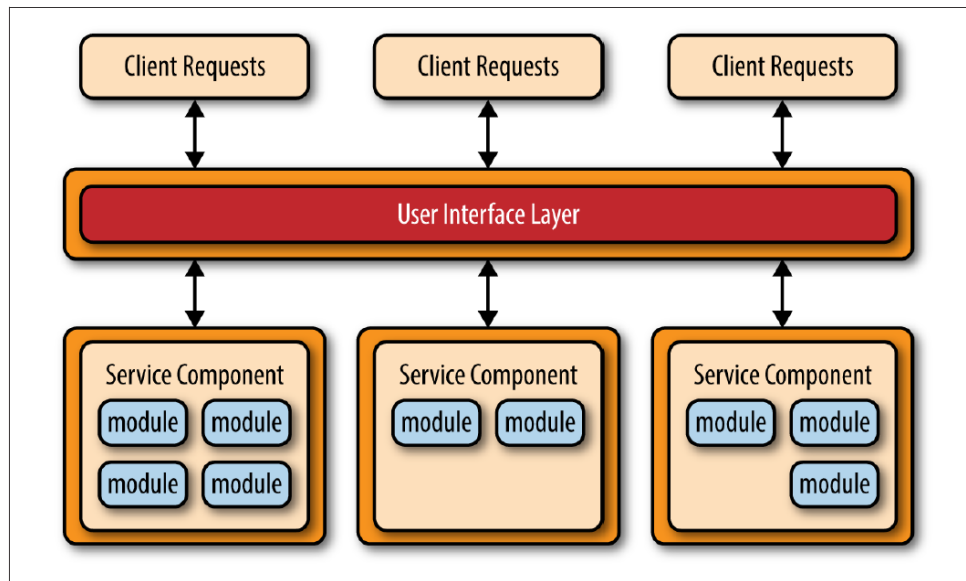


Figure 2 - Basic Micro-services architecture pattern

4.2 PATTERN TOPOLOGY

For JPO ODE, a *centralized messaging* topology is being envisioned. This topology (illustrated in Figure 3) uses a lightweight centralized message broker (e.g., Kafka). The lightweight message broker found in this topology does not perform any orchestration, transformation, or complex routing; rather, it is just a lightweight transport to access remote service components. The single point of failure and architectural bottleneck issues usually associated with a centralized broker are addressed through broker clustering.

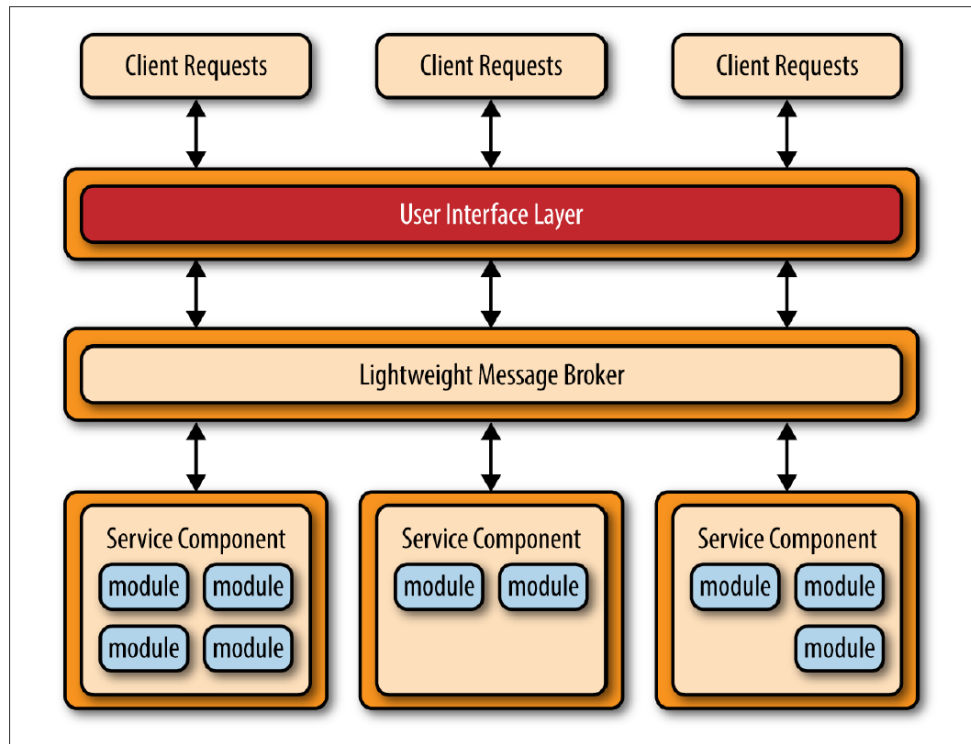


Figure 3 - Centralized messaging topology

Broker clustering refers to the ability of the message broker to scale horizontally and proportionally with the demands of the connected applications and services, ensuring the reliability of the messages rerouted through the broker. If needed, the message brokers can be distributed across multiple nodes to continue to provide services despite outages of one or more nodes and be able to scale in and out automatically as the data volume scales down and up.

If broker clustering is utilized, however, messages will not be guaranteed to be delivered in the same order as they arrived. In that case another caching service or data store will be responsible for re-ordering the messages based on a sequence key.

Apache Kafka is the messaging framework that will be incorporated in the JPO ODE implementation.

Figure 4 below highlights the concepts used in the Kafka implementation. Kafka has three key capabilities:

1. **Publish/Subscribe:** It lets you publish and subscribe to streams of records. In this respect it is similar to a message queue or enterprise messaging system.
2. **Persistent and Reliable:** It lets you store streams of records in a fault-tolerant way.
3. **Stream Processing:** It lets you process streams of records as they occur.

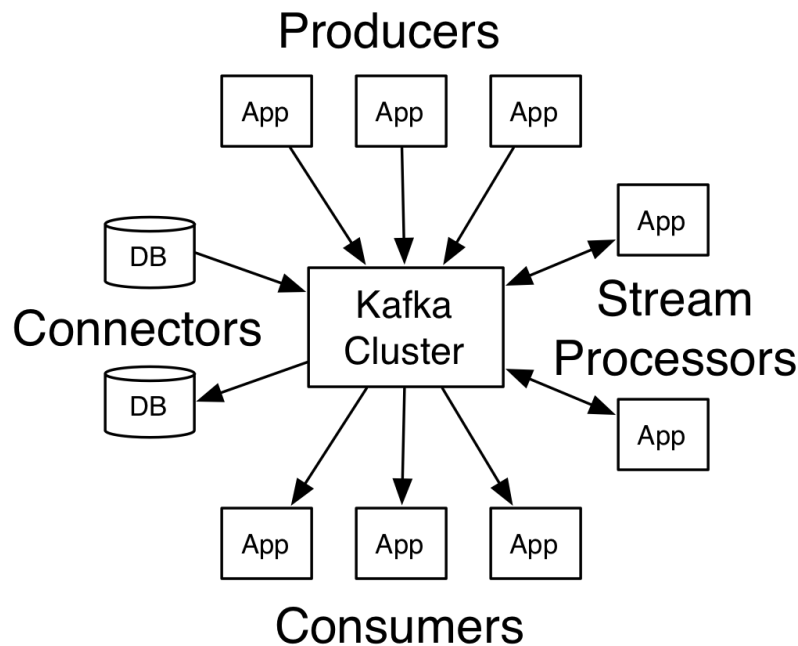


Figure 4 - Kafka Concepts

In order to connect to Kafka, there are 4 core API's that systems can use to communicate with the broker.

- The [Producer API](#) allows an application to publish a stream of records to one or more Kafka topics.
- The [Consumer API](#) allows an application to subscribe to one or more topics and receive a stream of records. Multiple applications can subscribe to a single topic and process messages in parallel via Kafka's consumer group handling.
- The [Streams API](#) allows an application to act as a *stream processor*, consuming an input stream from one or more topics and producing an output stream to one or more output topics, effectively transforming the input streams to output streams.
- The [Connector API](#) allows building and running reusable producers or consumers that connect Kafka topics to existing applications or data systems. For example, a connector to a relational database might capture every change to a table.

The ODE utilizes these Kafka concepts and the framework has been designed as depicted in Figure 5.

- The input services represent the publisher into the system.
- The BSM decoder service consumes an encoded topic and published a decoded topic.
- Applications such as the management console and gateways consume the outputted decoded messages.

Publish-Subscription Model

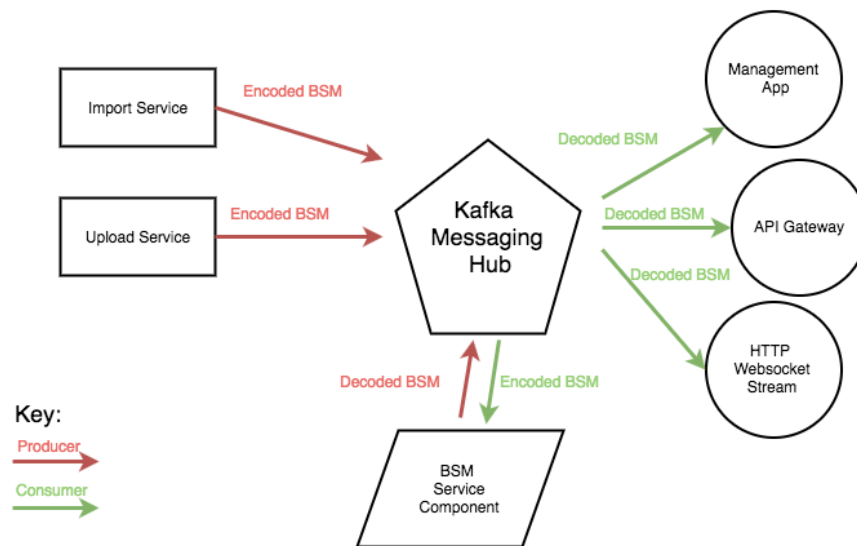


Figure 5 - Kafka Publish/Subscribe Model

5 JPO ODE MICRO-SERVICES TOPOLOGY

Figure 6 below represents the micro-services topology envisioned for JPO ODE. It highlights the granularity of ODE micro-services and identifies the major architectural components with which these services interact.

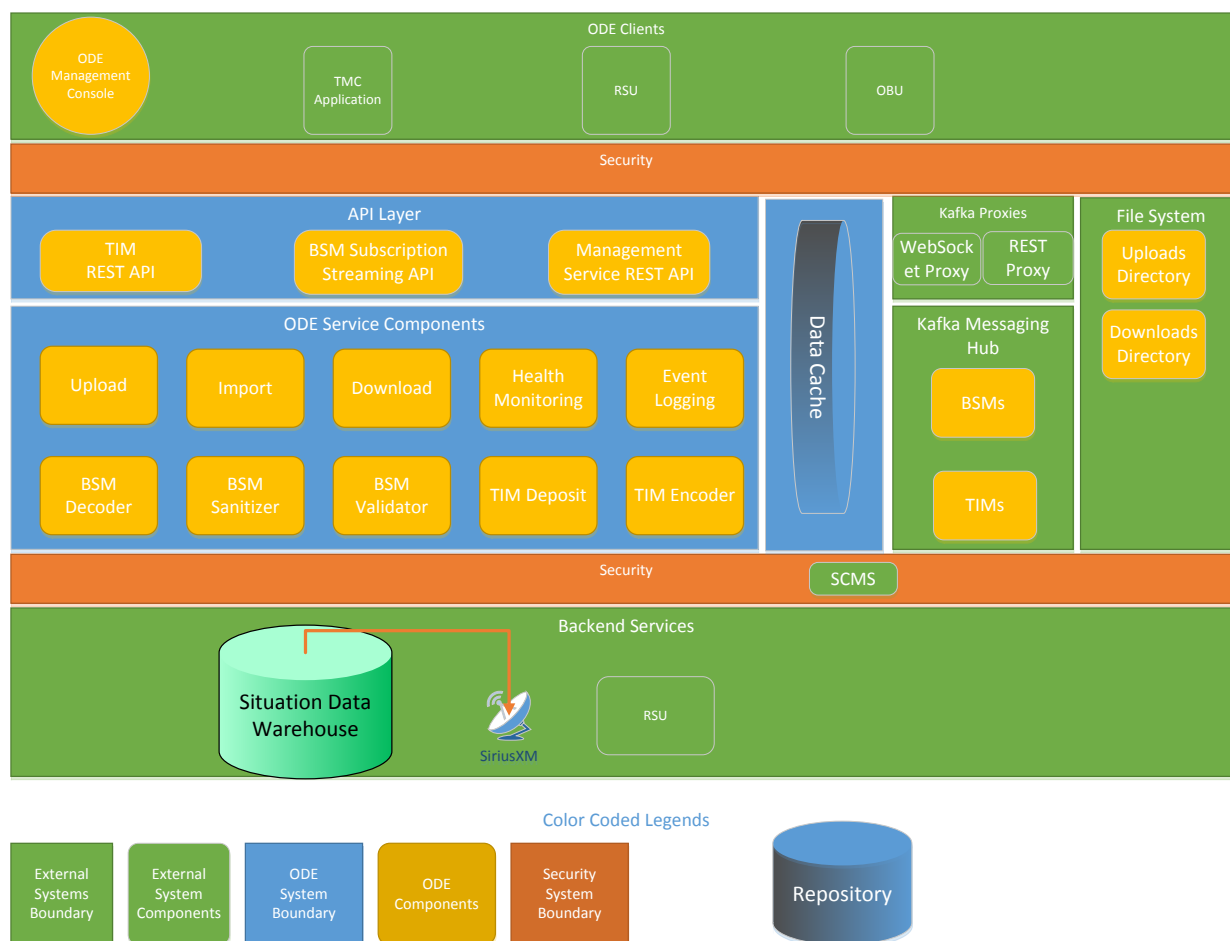


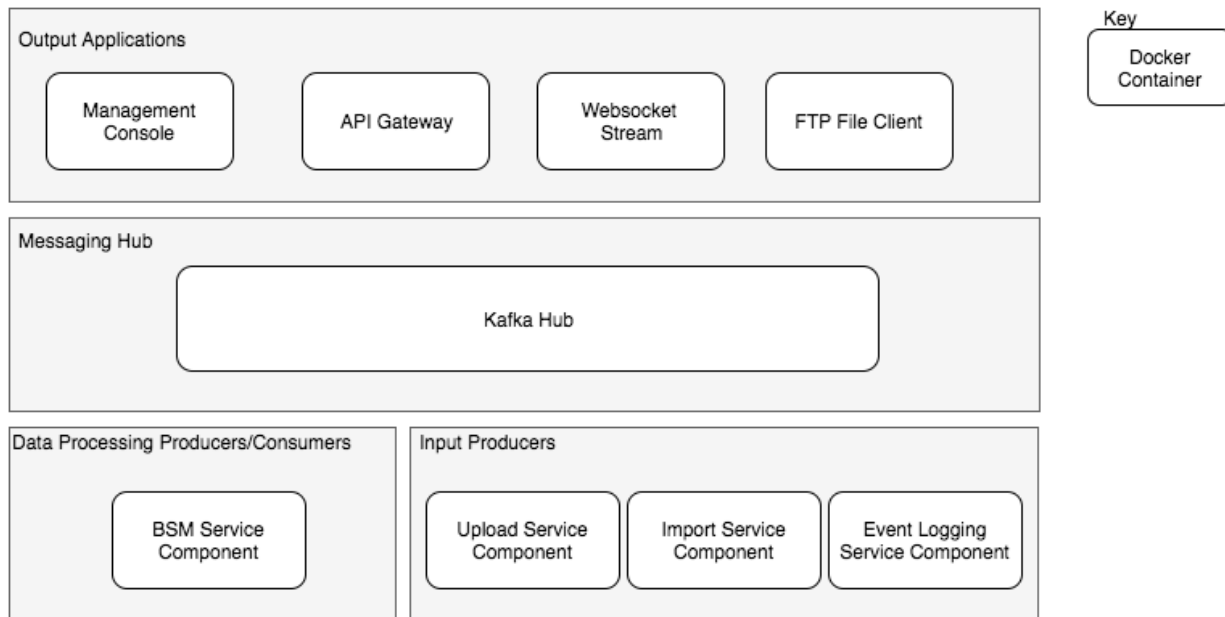
Figure 6 - JPO ODE Micro-services Topology

5.1 DEPLOYMENTS

Docker will be utilized as part of the deployment to compartmentalize each of the designed micro-services into separate containers. For the initial release(s) of the ODE, we will most likely be deploying the ODE as a semi-monolithic application where multiple services are provided to ODE clients as separate end-points rather than separate micro-services. We will utilize Docker to package up all of the components in a composite of containers each running some services. The ODE application itself will be running in one container and other major frameworks such as ZooKeeper and Kafka will be running in their own separate container(s). The ODE application will be the build artifact of the JPO ODE Continuous Integration and Delivery platform. As new services are developed, we will have the option of integrating them into the ODE monolithic application or as separate services that run in separate containers. The decision to make the service available as a standalone service will be made based on the needs of the customers and users of the JPO ODE. Figure 7 illustrates the two deployment options.

Docker Deployment Options

Full Micro-service Deployment



Monolithic Service Deployment

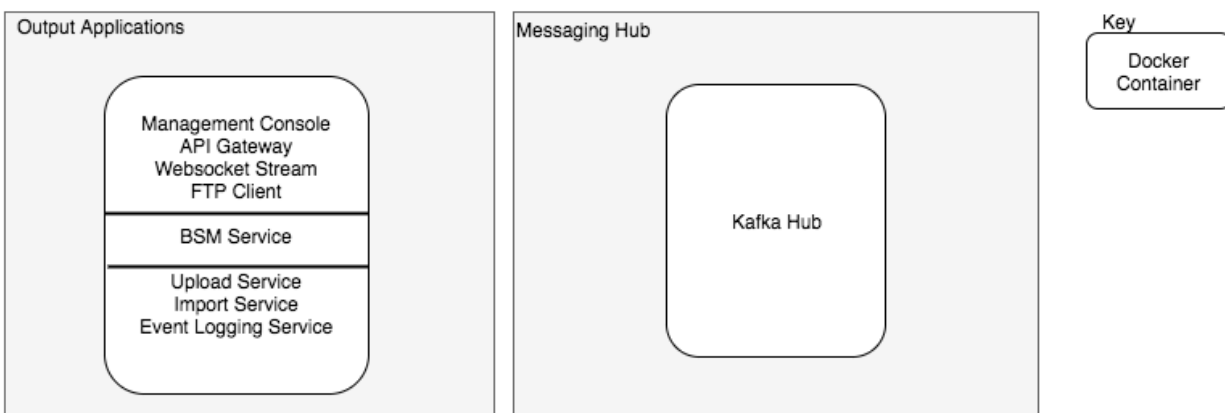


Figure 7 - Docker Deployment Options

6 APPENDIX

6.1 GLOSSARY

Term	Description
------	-------------

API	Application Program Interface
ASN.1	Abstract Syntax Notation One (ASN.1) is a standard and notation that describes rules and structures for representing, encoding, transmitting, and decoding data in telecommunications and computer networking
JPO	Joint Program Office
Kafka	Apache Kafka is publish-subscribe messaging rethought as a distributed commit log.
SCP	Secure Copy
US DOT	Unites States Department of Transportation
WebSocket	WebSocket is designed to be implemented in web browsers and web servers, but it can be used by any client or server application. The WebSocket Protocol is an independent TCP-based protocol. Its only relationship to HTTP is that its handshake is interpreted by HTTP servers as an Upgrade request.
ZooKeeper	Apache ZooKeeper is a centralized service for maintaining configuration information, naming, providing distributed synchronization, and providing group services.